



Core OS



About Me
CTO/CO-FOUNDER
systems engineer

@brandonphilips
github.com/philips



About Me
CTO/CO-FOUNDER
systems engineer

@brandonphilips
github.com/philips

Why build CoreOS?

The Datacenter as a Computer
*An Introduction to the Design of
Warehouse-Scale Machines*

Luiz André Barroso and Urs Hölzle
Google Inc.

containers

run and isolate apps

containers

what is it exactly?

libc

python

django

app.py

```
$ /usr/bin/python run app.py
```


libc
python
django
app.py

example.com/myapp

libc

python

django

app.py

```
$ container fetch example.com/myapp
```

```
$ container run example.com/myapp
```

pid ns

isolated pid 1

user ns

isolated uid 0

network ns

isolated netdev

mount ns

isolated /

cgroups

manage resources

cgroups

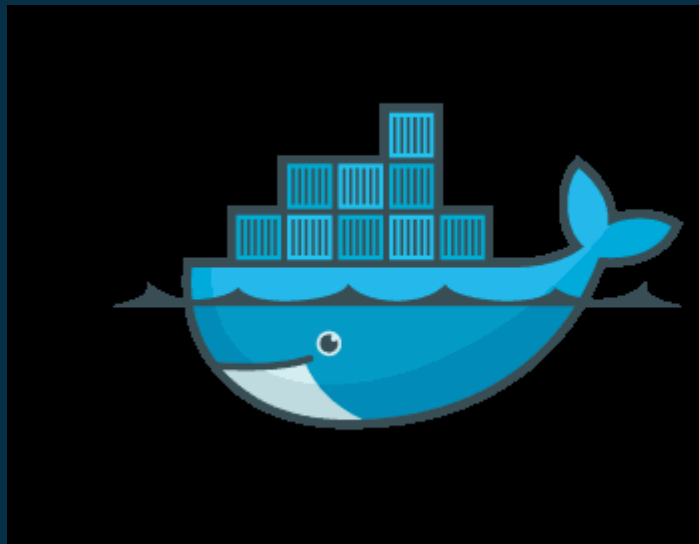
count resources

cAdvisor



cgroups

limit resources



docker engine



google lmctfy
cloud foundry garden
mesos containers

lxc
systemd-nspawn

containers


how are they created?

 Commit



 Github



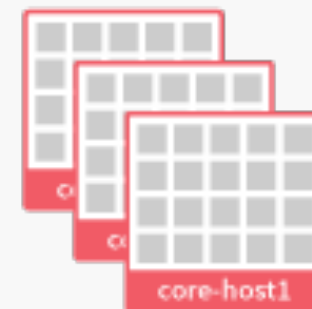
 CI System



 Registry



Docker Pull



containers

super-powers

App independence from the OS.

System to get container to the server.

Resource isolation between apps.

OS

reduced API contracts

kernel
systemd
etcd
ssh
docker

python
java
nginx
mysql
openssl

o distro distro distro distro distro distro

app

kernel
systemd
etc
ssh
docker

o distro distro distro distro distro distro

python
java
nginx
mysql
openssl

app

kernel
systemd
etc
ssh
docker

o distro distro distro distro distro distro

python
openssl-A

app1

java
openssl-B

app2

java
openssl-B

app3

manual
updates



automatic updates



automatic updates

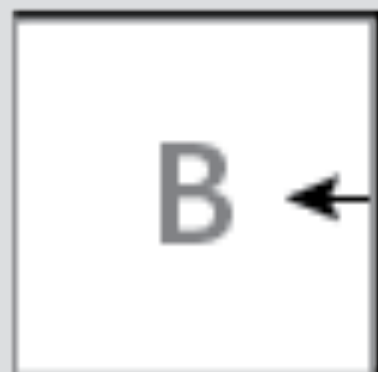
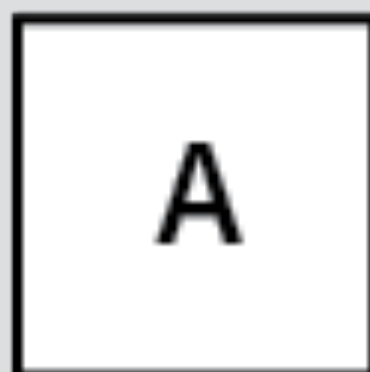


auto updates

atomic with rollback



Update



Data

A

B



OS

super-powers

Opportunity for automatic updates.

Consistent set of software across hosts.

Base OS independent from app.

clustering

design for host failure

etcd

/etc

distributed

open source software

sequentially consistent

exposed via HTTP

runtime reconfigurable

-X GET

Get Wait

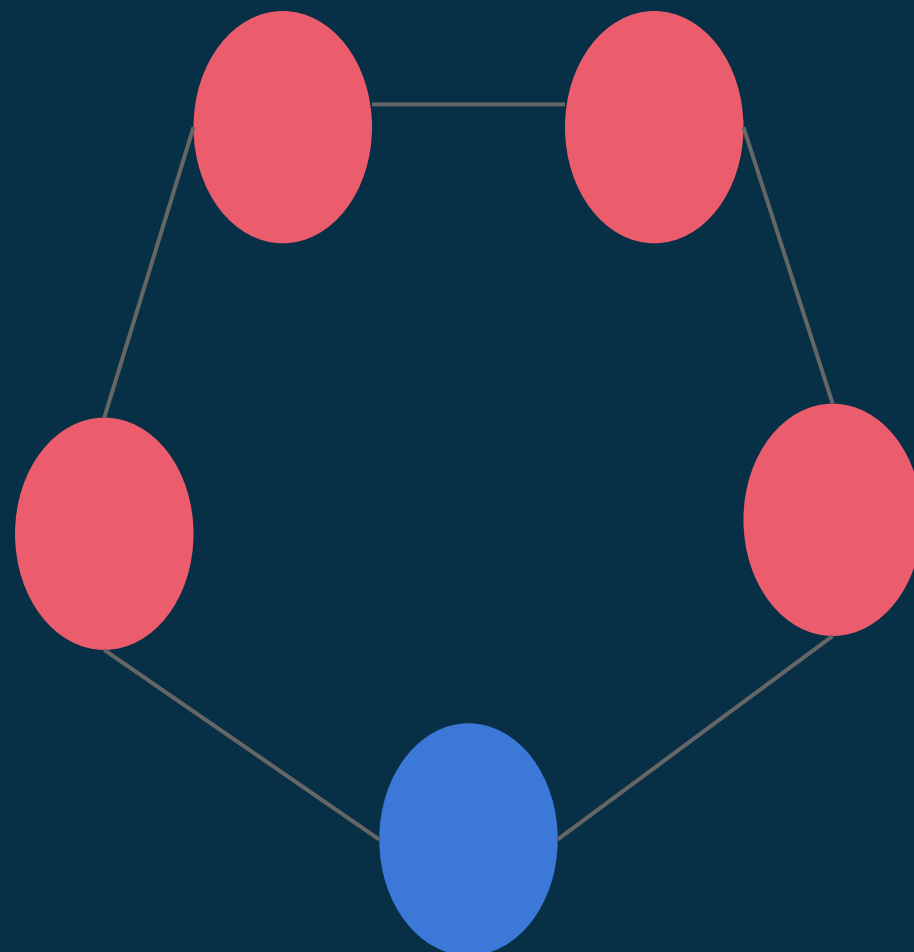
-X PUT

Put Create CAS

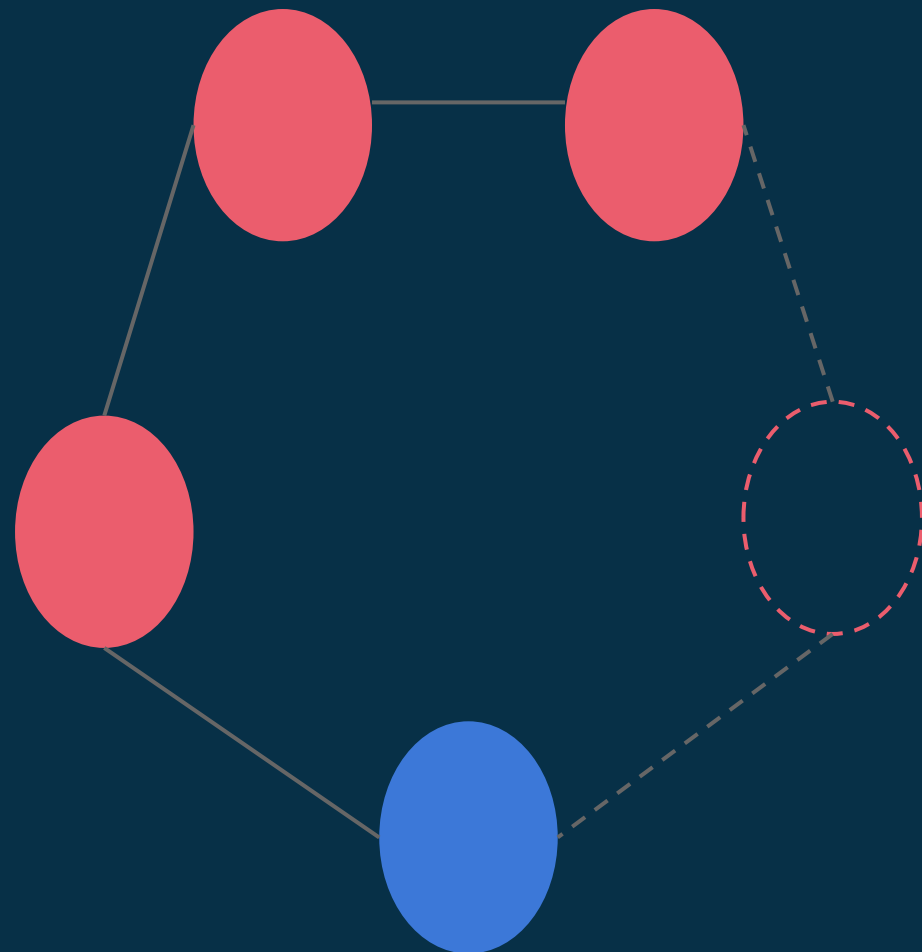
-X DELETE

Delete CAD

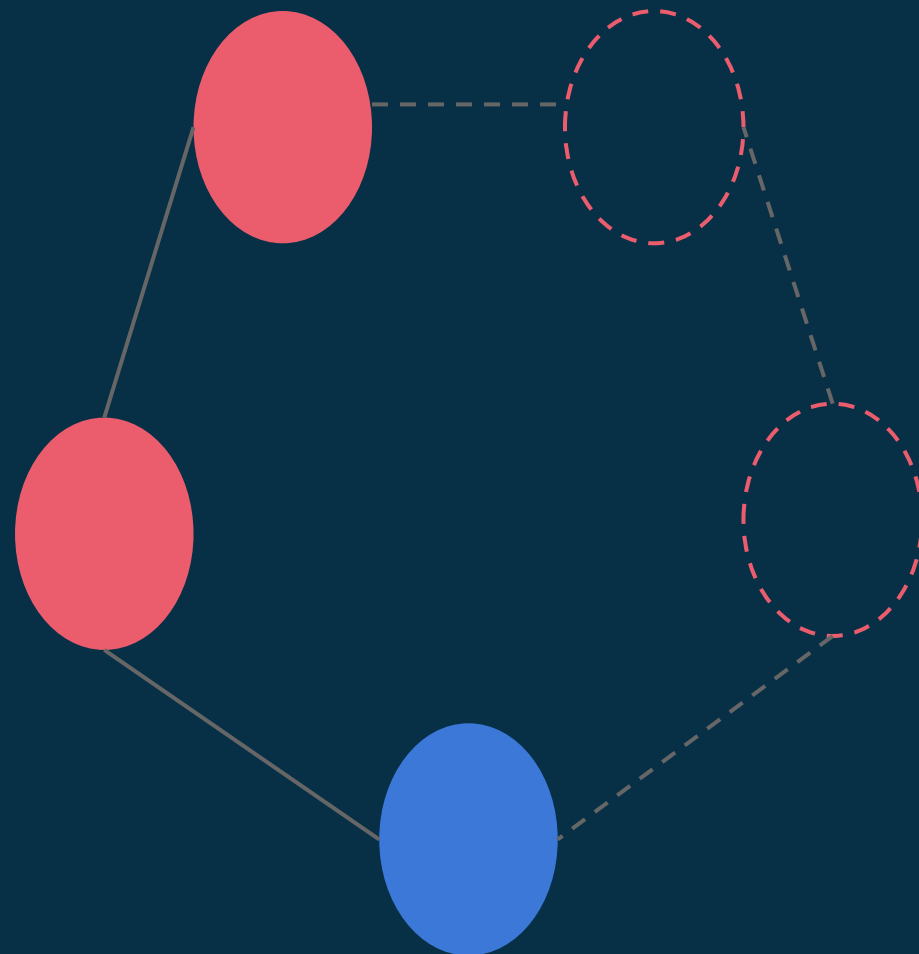
Available



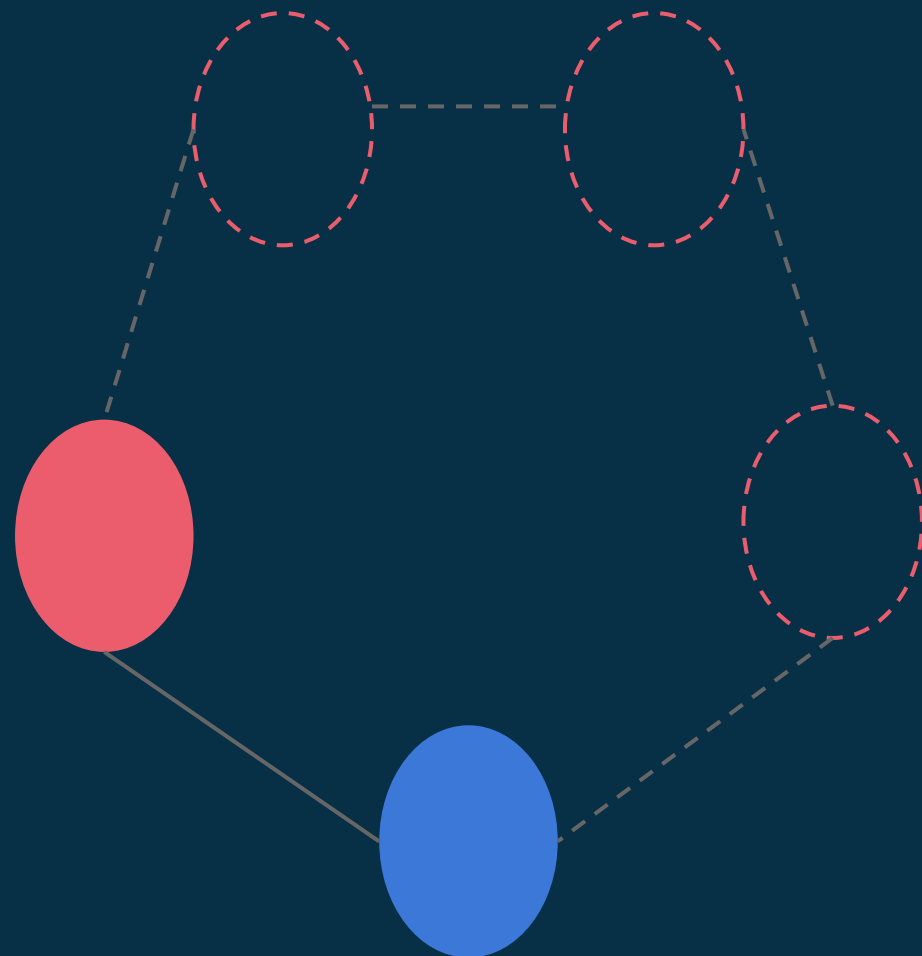
Available



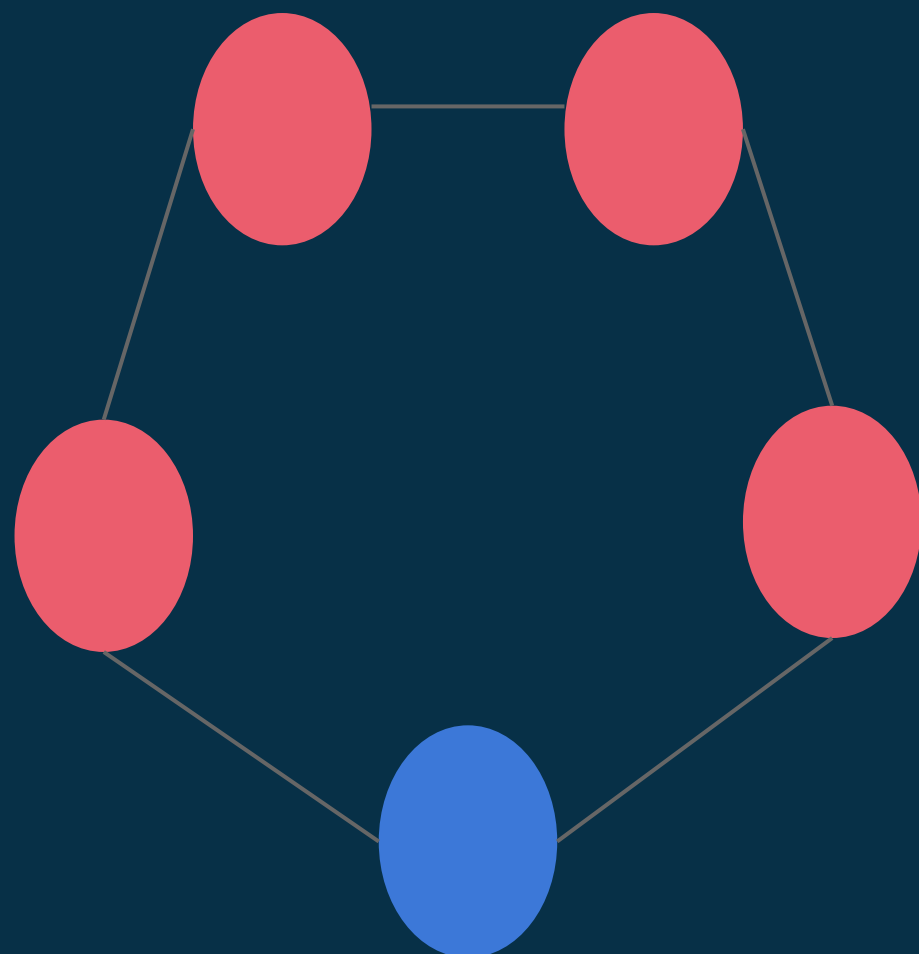
Available



Unavailable



Available

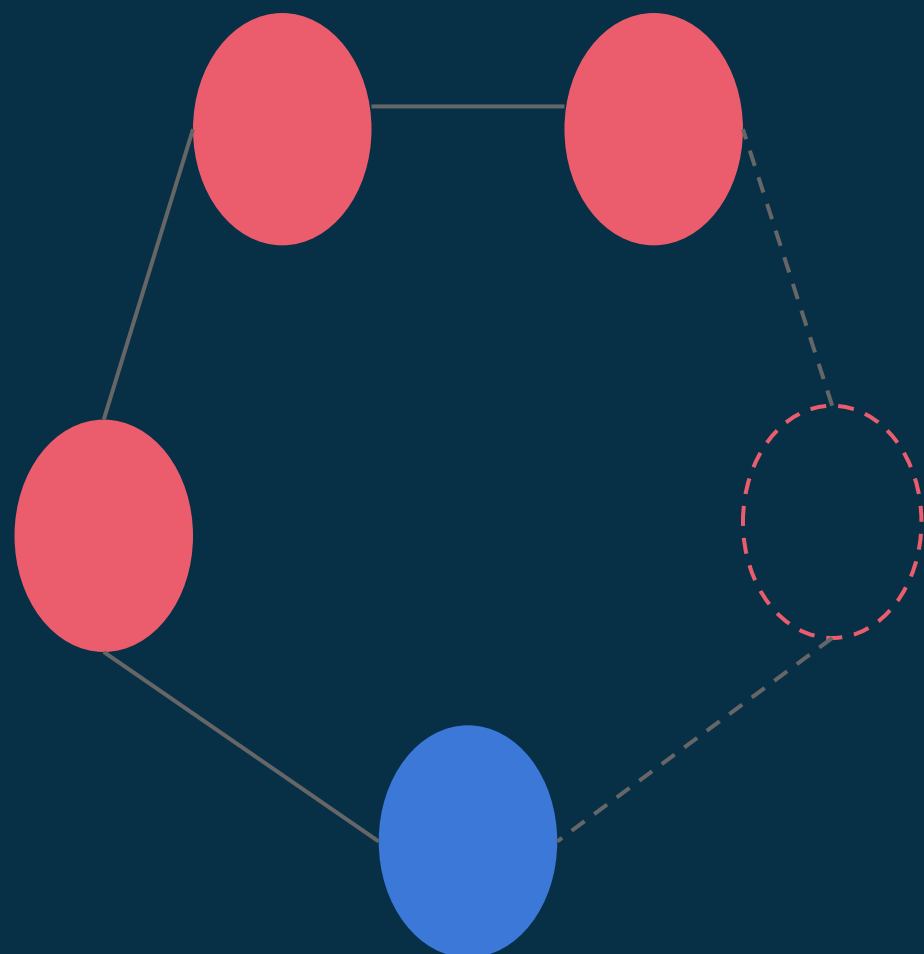


Leader

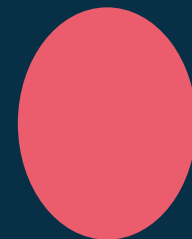


Follower

Available

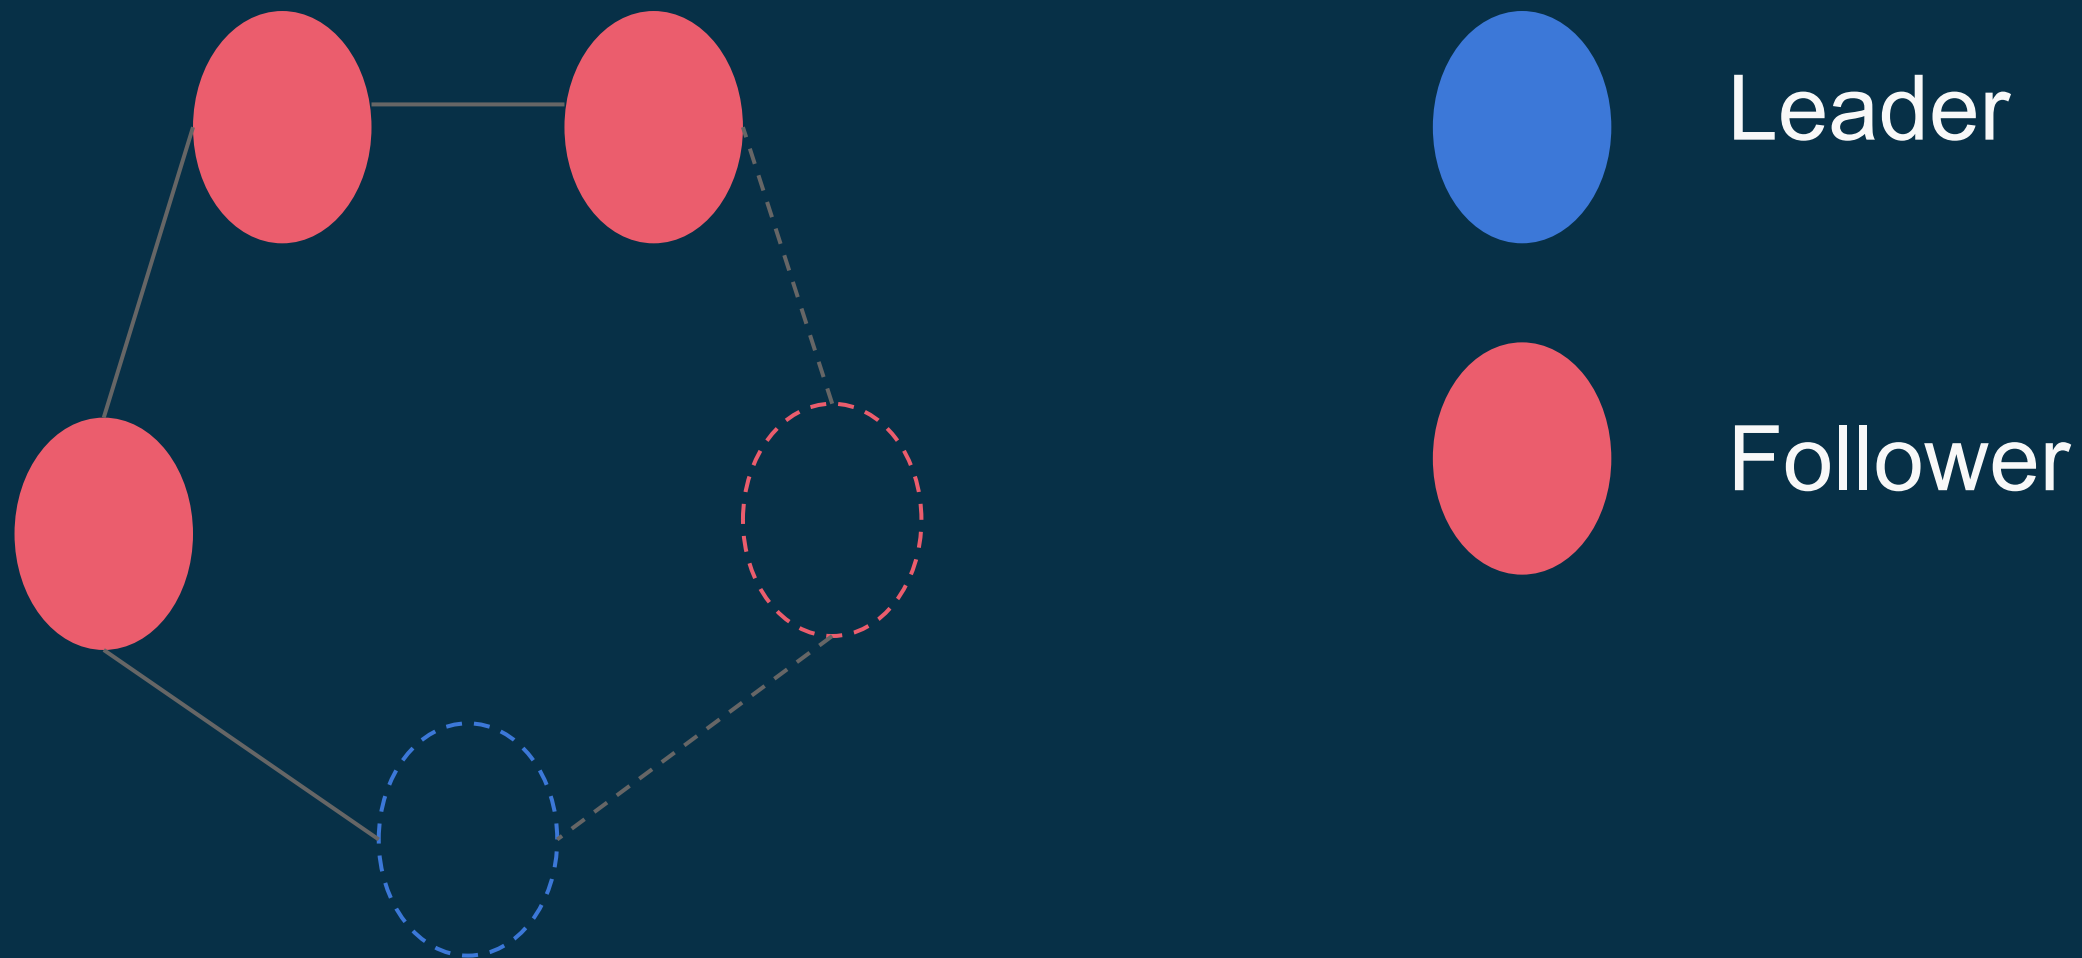


Leader

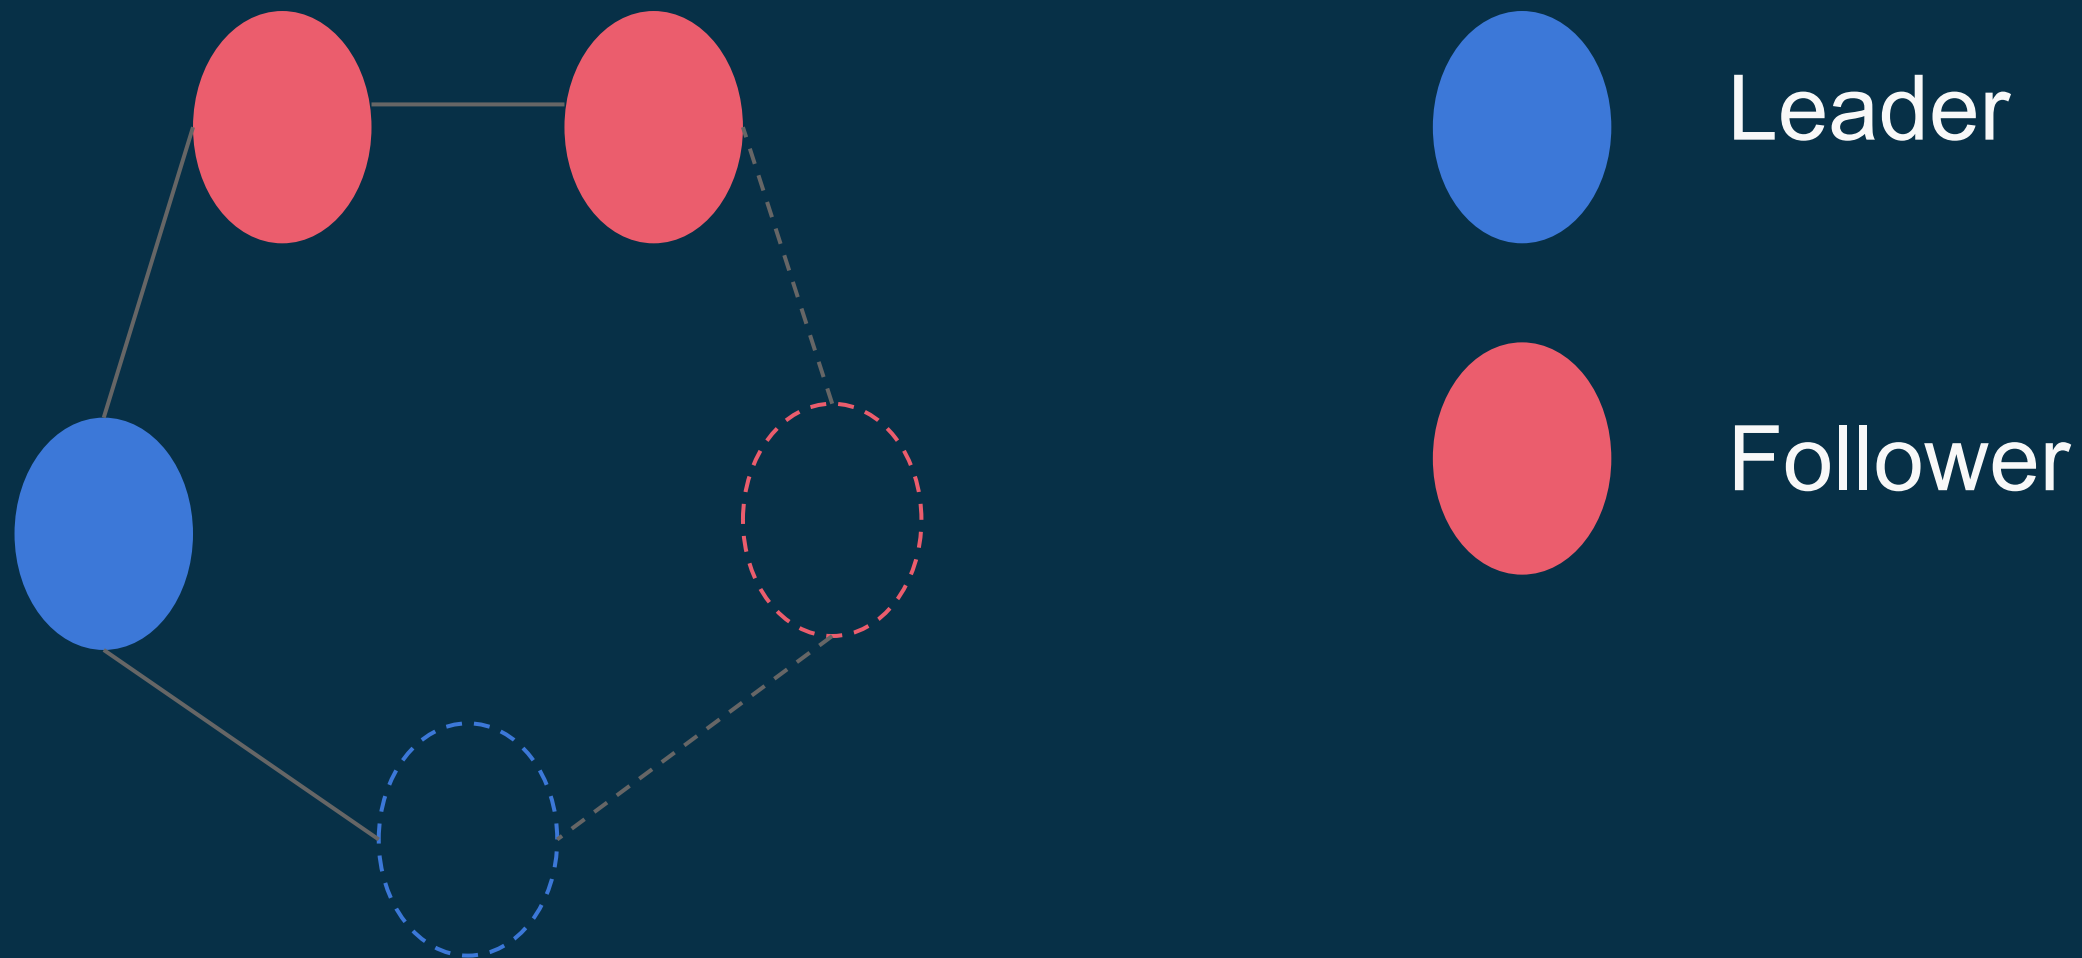


Follower

Temporarily Unavailable



Available



etcd

super-powers

Share configuration data across hosts.

Resilient to host failures.

Designed for consistency across hosts.

scheduling

getting work to servers

You



Scheduler API



Scheduler



Machine(s)

```
$ cat foo.service  
[Service]  
ExecStart=/usr/bin/sleep 500
```

```
$ fleetctl start foo.service  
Job foo.service launched on  
e1cd2bcd.../172.17.8.101
```



```
while true {  
  todo = diff(desState, curState)  
  schedule(todo)  
}
```

```
while true {  
  todo = diff(desState, curState)  
  schedule(todo)  
}
```

```
while true {  
  todo = diff(desState, curState)  
  schedule(todo)  
}
```

```
while true {  
  todo = diff(desState, curState)  
  schedule(todo)  
}
```

job scheduling

fleet

mesos

kubernetes

swarm

coordination

locksmith

scheduling

super-powers

Think about app capacity first.

Take advantage of compute resources.

Build for resilience to host failure.

service discovery

skydns, discoverd, confd

service discovery

magic proxies

OS

Containers

Cluster Configuration

Job Scheduling

Service Discovery



Core OS



Core OS

Go Beavs!